

Problem 1: Discrete LSPI (Scalar Basis)

Setup:

- State Space: $\mathcal{S} = \{-1, 1, 2\}$
- Action Space: $\mathcal{A} = \{-1, 0, 1\}$
- Transitions D : $\{(s_0 = 1, a_0 = 1, r_1 = 1, s_1 = 2), (2, 0, -1, 1), (1, -1, 0, -1)\}$
- Basis Function: $\phi(s, a) = a^2s + as + a$
- Initial Weight: $\omega^0 = 1$
- Discount: $\gamma = 0.9$

1.1 Derivation of Policy π^1

The policy is greedy with respect to $Q(s, a) = \phi(s, a)\omega^0$.

$$\pi^1(s) = \operatorname{argmax}_{a \in \mathcal{A}} (a^2s + as + a) \cdot 1$$

State $s = 2$:

$$\pi^1(2) = \operatorname{argmax}_a (2a^2 + 3a) = \operatorname{argmax}\{-1, 0, 5\} = 1$$

State $s = 1$:

$$\pi^1(1) = \operatorname{argmax}_a (a^2 + 2a) = \operatorname{argmax}\{-1, 0, 3\} = 1$$

State $s = -1$:

$$\pi^1(-1) = \operatorname{argmax}_a (-a^2) = \operatorname{argmax}\{-1, 0, -1\} = 0$$

Resulting Policy:

$$\pi^1(s) = \begin{cases} 0 & \text{if } s = -1 \\ 1 & \text{if } s = 1 \\ 1 & \text{if } s = 2 \end{cases}$$

1.2 Policy Evaluation (Calculation of ω^1)

We compute the scalar A and b using the LSTDQ statistics:

$$A = \frac{1}{L} \sum_{i=0}^{L-1} \phi(s_i, a_i) (\phi(s_i, a_i) - \gamma \phi(s_{i+1}, \pi^1(s_{i+1})))$$

$$b = \frac{1}{L} \sum_{i=0}^{L-1} \phi(s_i, a_i) r_{i+1}$$

Substituting the samples from D :

$$A = \frac{1}{3} \left[3(3 - 0.9(5)) + 0(0 - 0.9(3)) + (-1)(-1 - 0.9(0)) \right]$$

$$= \frac{1}{3} \left[3(-1.5) + 0 + 1 \right] = \frac{1}{3}[-3.5] = -1.1667$$

$$b = \frac{1}{3} \left[3(1) + 0(-1) + (-1)(0) \right] = \frac{1}{3}[3] = 1$$

Solving for weights:

$$\omega^1 = A^{-1}b = \frac{1}{-1.1667} = -0.8571$$

1.3 Derivation of Policy π^2

Update policy greedy w.r.t $Q(s, a) = \phi(s, a)(-0.8571)$.

$$\pi^2(s) = \operatorname{argmax}_{a \in \mathcal{A}} (a^2 s + a s + a)(-0.8571)$$

State $s = 2$:

$$\pi^2(2) = \operatorname{argmax}\{0.8571, 0, -4.2855\} = -1$$

State $s = 1$:

$$\pi^2(1) = \operatorname{argmax}\{0.8571, 0, -2.5713\} = -1$$

State $s = -1$:

$$\pi^2(-1) = \operatorname{argmax}\{0.8571, 0, 0.8571\} \rightarrow -1 \quad (\text{Tie-break})$$

Resulting Policy:

$$\pi^2(s) = -1 \quad \forall s \in \mathcal{S}$$

1.4 Policy Evaluation (Calculation of ω^2)

Next actions are all -1 .

$$\begin{aligned} A &= \frac{1}{3} \left[3(3 - 0.9(-1)) + 0(0 - 0.9(-1)) + (-1)(-1 - 0.9(-1)) \right] \\ &= \frac{1}{3} \left[3(3.9) + 0 + (-1)(-0.1) \right] = \frac{1}{3}[11.7 + 0.1] = 3.9333 \end{aligned}$$

$$b = \frac{1}{3} \left[3(1) + 0(-1) + (-1)(0) \right] = 1$$

Solving for weights:

$$\omega^2 = A^{-1}b = \frac{1}{3.9333} = 0.2542$$

1.5 Derivation of Policy π^3

Update policy greedy w.r.t $Q(s, a) = \phi(s, a)(0.2542)$.

$$\pi^3(2) = \operatorname{argmax}\{-0.2542, 0, 1.271\} = 1$$

$$\pi^3(1) = \operatorname{argmax}\{-0.2542, 0, 0.7626\} = 1$$

$$\pi^3(-1) = \operatorname{argmax}\{-0.2542, 0, -0.2542\} = 0$$

Conclusion: $\pi^3 = \pi^1$. The policy has converged.

Problem 2: Discrete LSPI (Vector Basis)

Setup:

- Same Data D .
- Basis Function: $\phi(s, a) = \begin{bmatrix} as + a \\ a^2 s \end{bmatrix}$.
- Initial Weights: $\omega^0 = [1, 1]^\top$.

2.1 Derivation of Policy π^1

$Q(s, a) = \phi(s, a)^\top \omega^0 = (as + a) + a^2 s$. This functional form is identical to Problem 1.

$$\pi^1(s) = \{1, 1, 0\} \quad \text{for } s = \{2, 1, -1\}.$$

2.2 Policy Evaluation (Calculation of ω^1)

We compute the matrix \mathbf{A} and vector \mathbf{b} .

$$\mathbf{A} = \frac{1}{L} \sum_{i=0}^{L-1} \phi_i (\phi_i - \gamma \phi'_i)^\top, \quad \mathbf{b} = \frac{1}{L} \sum_{i=0}^{L-1} \phi_i r_{i+1}$$

Term 1: $s = 1, a = 1 \rightarrow s' = 2, \pi(2) = 1$.

$$\phi = \begin{bmatrix} 2 \\ 1 \end{bmatrix}, \quad \phi' = \begin{bmatrix} 3 \\ 2 \end{bmatrix}. \quad \text{Diff} = \begin{bmatrix} -0.7 \\ -0.8 \end{bmatrix}.$$

$$\phi(\text{Diff})^\top = \begin{bmatrix} -1.4 & -1.6 \\ -0.7 & -0.8 \end{bmatrix}.$$

Term 2: $s = 2, a = 0 \rightarrow s' = 1, \pi(1) = 1$.

$$\phi = \begin{bmatrix} 0 \\ 0 \end{bmatrix}. \quad \text{Contribution is } \mathbf{0}.$$

Term 3: $s = 1, a = -1 \rightarrow s' = -1, \pi(-1) = 0$.

$$\phi = \begin{bmatrix} -2 \\ 1 \end{bmatrix}, \quad \phi' = \begin{bmatrix} 0 \\ 0 \end{bmatrix}. \quad \text{Diff} = \begin{bmatrix} -2 \\ 1 \end{bmatrix}.$$

$$\phi(\text{Diff})^\top = \begin{bmatrix} 4 & -2 \\ -2 & 1 \end{bmatrix}.$$

Aggregation:

$$\begin{aligned} \mathbf{A} &= \frac{1}{3} \left(\begin{bmatrix} -1.4 & -1.6 \\ -0.7 & -0.8 \end{bmatrix} + \begin{bmatrix} 4 & -2 \\ -2 & 1 \end{bmatrix} \right) = \begin{bmatrix} 0.867 & -1.2 \\ -0.9 & 0.067 \end{bmatrix}. \\ \mathbf{b} &= \frac{1}{3} \left(\begin{bmatrix} 2 \\ 1 \end{bmatrix} (1) + \mathbf{0} + \begin{bmatrix} -2 \\ 1 \end{bmatrix} (0) \right) = \begin{bmatrix} 0.67 \\ 0.33 \end{bmatrix}. \end{aligned}$$

Solving $\mathbf{A}\omega^1 = \mathbf{b}$:

$$\omega^1 = \mathbf{A}^{-1} \mathbf{b} \approx \begin{bmatrix} -0.43 \\ -0.87 \end{bmatrix}.$$

2.3 Derivation of Policy π^2

Maximize $Q(s, a) = -0.43(as + a) - 0.87(a^2s)$. For $s = 1$, $Q(1, 0) = 0$ and $Q(1, -1) \approx 0$. The tie is broken with -1 .

Resulting Policy: $\pi^2 = \{0, -1, -1\}$.

2.4 Policy Evaluation (Calculation of ω^2)

Next actions: $\pi^2(2) = 0, \pi^2(1) = -1, \pi^2(-1) = -1$.

Aggregation: Note: For the third term, $\pi(-1) = -1 \implies a = -1, s = -1$. Thus $\phi' = [as + a, a^2s]^\top = [0, -1]^\top$.

$$\begin{aligned}\mathbf{A} &= \frac{1}{3} \left[\begin{bmatrix} 2 \\ 1 \end{bmatrix} \left(\begin{bmatrix} 2 \\ 1 \end{bmatrix} - 0.9\mathbf{0} \right)^\top + \mathbf{0} + \begin{bmatrix} -2 \\ 1 \end{bmatrix} \left(\begin{bmatrix} -2 \\ 1 \end{bmatrix} - 0.9 \begin{bmatrix} 0 \\ -1 \end{bmatrix} \right)^\top \right] \\ &= \begin{bmatrix} 2.67 & -0.6 \\ 0 & 0.967 \end{bmatrix}. \\ \mathbf{b} &= \begin{bmatrix} 0.67 \\ 0.33 \end{bmatrix}.\end{aligned}$$

Solving $\mathbf{A}\omega^2 = \mathbf{b}$:

$$\omega^2 \approx \begin{bmatrix} 0.33 \\ 0.34 \end{bmatrix}.$$

2.5 Derivation of Policy π^3

Maximize $Q(s, a) = 0.33(as + a) + 0.34(a^2s)$. Since weights are positive, the logic mirrors π^1 .

Resulting Policy: $\pi^3 = \{1, 1, 0\}$. This matches π^1 .

Problem 3: Continuous State LSPI

Setup:

- Continuous States $\mathcal{S} \in [-2, 2]$. Continuous Actions $\mathcal{A} \in [-1.5, 2]$.
- Data: $D = \{(2, 1, 1, 1), (1, -1, 2, -1)\}$.
- Basis: $\phi(s, a) = [s, sa]^\top$. $\omega^0 = [0.5, 1]^\top$.

3.1 Derivation of Policy π^1

Maximize $Q(s, a) = 0.5s + 1(sa) = s(0.5 + a)$ with respect to $a \in [-1.5, 2]$.

- If $s > 0$: Maximize $(0.5 + a) \implies a = 2$.
- If $s < 0$: Minimize $(0.5 + a)$ to make product positive $\implies a = -1.5$.

For Data points: $\pi^1(1) = 2, \pi^1(-1) = -1.5$.

3.2 Policy Evaluation (Calculation of ω^1)

Aggregation:

$$\mathbf{A} = \frac{1}{2} \begin{bmatrix} 4.1 & -1.95 \\ 0.3 & 2.75 \end{bmatrix} = \begin{bmatrix} 2.05 & -0.975 \\ 0.15 & 1.375 \end{bmatrix}.$$

$$\mathbf{b} = \frac{1}{2} ([2, 2]^\top(1) + [1, -1]^\top(2)) = \begin{bmatrix} 2 \\ 0 \end{bmatrix}.$$

Solving $\mathbf{A}\omega^1 = \mathbf{b}$ yields $\omega^1 \approx [0.93, -0.10]^\top$.

3.3 Derivation of Policy π^2

Maximize $Q(s, a) = 0.93s - 0.10sa = s(0.93 - 0.10a)$.

- If $s > 0$: Maximize $(0.93 - 0.10a) \implies \text{Min } a \implies a = -1.5$.
- If $s < 0$: Minimize $(0.93 - 0.10a) \implies \text{Max } a \implies a = 2$.

For Data points: $\pi^2(1) = -1.5$, $\pi^2(-1) = 2$.

3.4 Policy Evaluation (Calculation of ω^2)

Aggregation: Note: For the second term, $\pi(-1) = 2 \implies sa = (-1)(2) = -2$. Thus $\phi' = [-1, -2]^\top$.

$$\mathbf{A} = \frac{1}{2} \left(\begin{bmatrix} 2 \\ 2 \end{bmatrix} \left(\begin{bmatrix} 2 \\ 2 \end{bmatrix} - 0.9 \begin{bmatrix} 1 \\ -1.5 \end{bmatrix} \right)^\top + \begin{bmatrix} 1 \\ -1 \end{bmatrix} \left(\begin{bmatrix} 1 \\ -1 \end{bmatrix} - 0.9 \begin{bmatrix} -1 \\ -2 \end{bmatrix} \right)^\top \right)$$

$$= \begin{bmatrix} 2.05 & 3.75 \\ 0.15 & 2.95 \end{bmatrix}.$$

$$\mathbf{b} = \begin{bmatrix} 2 \\ 0 \end{bmatrix}.$$

Solving for weights:

$$\omega^2 \approx \begin{bmatrix} 1.076 \\ -0.055 \end{bmatrix}.$$

3.5 Interpretation of Final Policy π^3

Maximize $Q(s, a) = s(1.076 - 0.055a)$. The logic remains the same as π^2 .

$$\pi^3(s) = \begin{cases} -1.5 & \text{if } s > 0 \\ 2 & \text{if } s < 0 \end{cases}$$